

1. Introduction to Dataset

2. Exploratory Data Analysis

- Checking for Missing Data.
- Class Imbalance
- Visualizing distributions.
- 'Amount' Distribution
- 'Time' Distribution

3. Modeling Outcome of Interest

- The Problem of Imbalanced Data (How NOT to do it...)
- Feature Engineering
 - Time-Based Features
 - Feature: Time_hour > 4
 - Feature: \$0 Fraud Amounts...?
 - Normalize Time and Amount
 - Add the Rest: PCA and Class
 - Classification Improvements after Feature Engineering
- Data Processing
 - Balancing Classes
 - Removing High-Correlation Outliers
 - Feature Selection
- Test and Compare Classifiers
 - First-Run: Predictions on Default Parameters
 - Logistic Regression- GridSearch & Recall Score.
- Pyrrhic Victory-
- Optimize Specificity, while Maintaining 100% Recall
 - Custom Scoring Function
 - Iteration Function
 - LogisticRegression- Optimized.
 - DecisionTreeClassifier- Optimized
 - Support Vector Classifier- Optimized
 - KNeighborsClassifier- Optimized
 - Imblearn' BalancedRandomForest- Optimized
 - SKlearn' RandomForestClassifier- Optimized

4. Research Question

5. Choosing Model

- Perfect Recall
- Best Overall

6. Practical Use for Audiences of Interest

7. Weak Points & Shortcomings

1. Introduction to Dataset

From <https://www.kaggle.com/mlg-ulb/creditcardfraud/home> :

The datasets contains transactions made by credit cards in September 2013 by european cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

PCA Features:

It contains only numerical input variables which are the result of a PCA transformation. Features V1, V2, ... V28 are the principal components obtained with PCA.

Time:

Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset.

Amount:

The feature 'Amount' is the transaction Amount.

Class:

Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import collections
6
7 # Classifier Libraries
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.svm import SVC
10 from sklearn.neighbors import KNeighborsClassifier
```

```
11 from sklearn.tree import DecisionTreeClassifier
12 from sklearn.ensemble import RandomForestClassifier
13 from imblearn.ensemble import BalancedRandomForestClassifier
14
15
16 # Other Libraries
17 from sklearn.model_selection import train_test_split, StratifiedShuffleSplit, GridSearchCV, cross_val_score
18 from sklearn.pipeline import make_pipeline
19 from imblearn.pipeline import make_pipeline as imbalanced_make_pipeline
20 from imblearn.over_sampling import SMOTE, ADASYN
21 from imblearn.under_sampling import RandomUnderSampler
22 from sklearn.metrics import make_scorer, precision_score, recall_score, classification_report, confusion_matrix
23 from collections import Counter
24 from sklearn.preprocessing import RobustScaler
25 import warnings
26 warnings.filterwarnings("ignore")
27
28
29 data = pd.read_csv('Data/creditcard.csv', sep=',')
30 data.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431

5 rows × 31 columns

2. Exploratory Data Analysis

```
1 print(data.columns)
```

```
1 Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
2        'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
3        'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
4        'Class'],
5        dtype='object')
```

```
1 data.shape
```

```
1 (284807, 31)
```

Checking for Missing Data.

- Fortunately, data integrity is perfect. All non-null, and no mixed types.

```
1 data.info()
```

```
1 <class 'pandas.core.frame.DataFrame'>
2 RangeIndex: 284807 entries, 0 to 284806
3 Data columns (total 31 columns):
4 Time      284807 non-null float64
5 V1        284807 non-null float64
6 V2        284807 non-null float64
7 V3        284807 non-null float64
8 V4        284807 non-null float64
9 V5        284807 non-null float64
10 V6        284807 non-null float64
11 V7        284807 non-null float64
12 V8        284807 non-null float64
13 V9        284807 non-null float64
14 V10       284807 non-null float64
15 V11       284807 non-null float64
16 V12       284807 non-null float64
```

```

17 V13      284807 non-null float64
18 V14      284807 non-null float64
19 V15      284807 non-null float64
20 V16      284807 non-null float64
21 V17      284807 non-null float64
22 V18      284807 non-null float64
23 V19      284807 non-null float64
24 V20      284807 non-null float64
25 V21      284807 non-null float64
26 V22      284807 non-null float64
27 V23      284807 non-null float64
28 V24      284807 non-null float64
29 V25      284807 non-null float64
30 V26      284807 non-null float64
31 V27      284807 non-null float64
32 V28      284807 non-null float64
33 Amount   284807 non-null float64
34 Class     284807 non-null int64
35 dtypes: float64(30), int64(1)
36 memory usage: 67.4 MB

```

Class Imbalance

- This is the most unique quality about this dataset. Most of the steps taken later will be about multiple ways of dealing with imbalanced data.

```

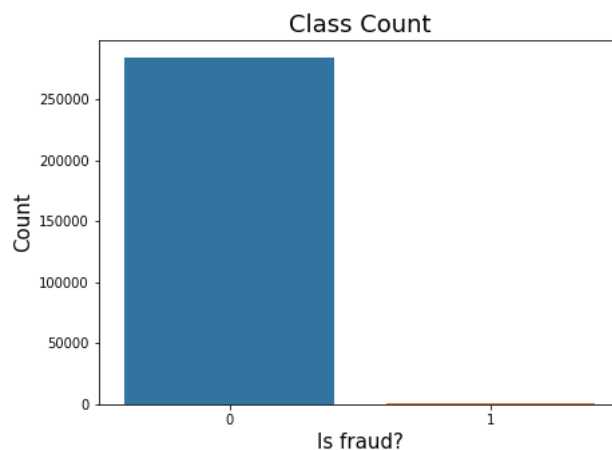
1 #Lets start looking the difference by Normal and Fraud transactions
2 print("Distribution of Normal(0) and Frauds(1): ")
3 print(data["Class"].value_counts())
4 print('')
5
6 # The classes are heavily skewed we need to solve this issue later.
7 print('Non-Frauds', round(data['Class'].value_counts()[0]/len(data) * 100,2), '% of the dataset')
8 print('Frauds', round(data['Class'].value_counts()[1]/len(data) * 100,2), '% of the dataset')
9
10 plt.figure(figsize=(7,5))
11 sns.countplot(data['Class'])
12 plt.title("Class Count", fontsize=18)
13 plt.xlabel("Is fraud?", fontsize=15)
14 plt.ylabel("Count", fontsize=15)
15 plt.show()

```

```

1 Distribution of Normal(0) and Frauds(1):
2 0      284315
3 1         492
4 Name: Class, dtype: int64
5
6 Non-Frauds 99.83 % of the dataset
7 Frauds 0.17 % of the dataset

```





Visualizing distributions.

- Features have different central tendencies and need to be normalized to make better sense of them.
- 'Time' is encoded in seconds, out of a 24Hr day. We'll need to transform it in order to visualize it properly.

```

1 plt.figure(figsize=(16,4))
2 data.iloc[:, :-1].boxplot()
3 plt.title('(Raw) Distribution of Features', fontsize=17)
4 plt.show()
5
6 plt.figure(figsize=(16,4))
7 np.log(data.iloc[:, :-1]).boxplot()
8 plt.title('(Log) Distribution of Features', fontsize=17)
9 plt.show()

```



- It's clear that `Time` and `Amount` are in a different range compared to the `PCA` features.

'Amount' Distribution

- Variable isn't normalized.
- There's high concentrations of small-amount transactions. And many dispersed large-amount outliers, all the way up to \$25,000
- 85\% of data is below \$140
- Top 1% of transaction amounts are between 1017.97 and 25691.16



Amount of frauds

- 80% of Frauds are less than: \$152.34.

```

1 #Now look at Fraud Amounts
2 plt.figure(figsize=(16,5))
3 sns.boxplot(x=data.Amount[data.Class == 1])
4 plt.title('Distribution of (Fraud) Amounts', fontsize=17)
5 plt.show()
6 #Now look at Non-Fraud Amounts
7 plt.figure(figsize=(16,5))
8 sns.boxplot(x=data.Amount[data.Class == 0])
9 plt.title('Distribution of (Non-Fraud) Amounts', fontsize=17)
10 plt.show()

```




```
1 print('Top 85% of transaction amounts:', round(data.Amount.quantile(.85),2))
2 print('Top 1% of transaction amounts:', round(data.Amount.quantile(.99),2))
3 print('Largest transaction amount:', round(data.Amount.quantile(1),2))
4 print('80% of Frauds are less than:', round(data.Amount[data.Class==1].quantile(.80),2))
```

```
1 Top 85% of transaction amounts: 140.0
2 Top 1% of transaction amounts: 1017.97
3 Largest transaction amount: 25691.16
4 80% of Frauds are less than: 152.34
```

'Time' Distribution

- I'll convert 'Time' to hours and minutes, which will allow for better visualization.
- 'Time' distribution (by second) shows two normal curves, which might reveal something meaningful for predicting purposes. This will be the basis for a time-based feature engineering.

```
1 #First look at Time
2 plt.figure(figsize=(11,6))
3 sns.distplot(data.Time,kde=False)
4 plt.title('Distribution of Time', fontsize=17)
5 plt.show()
```



```


1 # Create a EDA dataframe for the time units and visualizations
2 eda = pd.DataFrame(data.copy())
3
4 # Tell timedelta to interpret the Time as second units
5 timedelta = pd.to_timedelta(eda['Time'], unit='s')
6
7 # Create a hours feature from timedelta
8 eda['Time_hour'] = (timedelta.dt.components.hours).astype(int)

```

```

1 #Exploring the distribution by Class types through seconds
2 plt.figure(figsize=(12,5))
3 sns.distplot(eda[eda['Class'] == 0]["Time"],
4             color='g')
5 sns.distplot(eda[eda['Class'] == 1]["Time"],
6             color='r')
7 plt.title('(Density Histogram) Fraud VS Normal Transactions by Second', fontsize=17)
8 plt.xlim([-2000,175000])
9 plt.show()


```



```

1 #Exploring the distribution by Class types through hours
2 plt.figure(figsize=(12,5))
3 sns.distplot(eda[eda['Class'] == 0]["Time_hour"],
4             color='g')
5 sns.distplot(eda[eda['Class'] == 1]["Time_hour"],
6             color='r')
7 plt.title('(Density Histogram) Fraud VS Normal Transactions by Hour', fontsize=17)
8 plt.xlim([-1,25])
9 plt.show()

```



3. Modeling Outcome of Interest

The Problem of Imbalanced Data (How NOT to do it...)

- Here I'll do a base-line prediction of frauds using default settings on the data without any modifications.
- This serves to show the need for techniques on Class Imbalance.

Approach

- Below I split data into train and test groups.
- I'll make sure the groups maintain the same class balance as the whole set. That way they can better represent the whole, for testing purposes.

```

1 # Define outcome and predictors to split into train and test groups
2 y = data['Class']
3 X = data.drop('Class', 1)
4
5 X_train, X_test, y_train, y_test = train_test_split(
6     X, y, test_size=0.2, random_state=42)
7
8 # Class balance in test group
9 print("TEST GROUP")
10 print('Size:', y_test.count())
11 print("Frauds percentage:",
12       y_test.value_counts()[1]/y_test.count())
13 print("Nonfrauds percentage:",
14       y_test.value_counts()[0]/y_test.count())
15
16 # Class balance in train group
17 print("\nTRAIN GROUP")
18 print('Size:', y_train.count())
19 print("Frauds percentage:",
20       y_train.value_counts()[1]/y_train.count())
21 print("Nonfrauds percentage:",
22       y_train.value_counts()[0]/y_train.count())

```

```

1 TEST GROUP
2 Size: 56962
3 Frauds percentage: 0.0017204452090867595
4 Nonfrauds percentage: 0.9982795547909132
5
6 TRAIN GROUP
7 Size: 227845
8 Frauds percentage: 0.001729245759178389
9 Nonfrauds percentage: 0.9982707542408216

```

```

1 # Invoke classifier
2 clf = LogisticRegression()
3
4 # Cross-validate on the train data
5 train_cv = cross_val_score(X=X_train, y=y_train, estimator=clf, cv=3)
6 print("TRAIN GROUP")
7 print("\nCross-validation accuracy scores:", train_cv)
8 print("Mean score:", train_cv.mean())

```

```


9
10 # Now predict on the test group
11 print("\nTEST GROUP")
12 y_pred = clf.fit(X_train, y_train).predict(X_test)
13 print("\nAccuracy score:", clf.score(X_test, y_test))
14
15 # Classification report
16 print('\nClassification report:\n')
17 print(classification_report(y_test, y_pred))
18
19 # Confusion matrix
20 conf_matrix = confusion_matrix(y_test, y_pred)
21 sns.heatmap(conf_matrix, annot=True, fmt='d', cmap=plt.cm.copper)
22 plt.show()
23

```

```

1 TRAIN GROUP
2
3 Cross-validation accuracy scores: [0.99906516 0.99897298 0.99873598]
4 Mean score: 0.9989247069698471
5
6 TEST GROUP
7
8 Accuracy score: 0.9989993328885924
9
10 Classification report:
11
12             precision    recall  f1-score   support
13
14      0           1.00       1.00       1.00     56864
15      1           0.83       0.53       0.65        98
16
17   micro avg       1.00       1.00       1.00     56962
18   macro avg       0.91       0.77       0.82     56962
19   weighted avg     1.00       1.00       1.00     56962

```



Understanding the scores

Sensitivity (or Recall) is the percentage of positives correctly identified.

Specificity is just the opposite, the percentage of negatives correctly identified.

The confusion matrix and classification reports reveal that **the high scores are merely a reflection of the class imbalance**. Since we're using a generalized scoring method, accuracy reflects the recall of both frauds and non-frauds. However, since frauds are so few, (0.0017%) their poor recall (53%) isn't reflected in the overall accuracy score.

On the test set

- Of 98 fraud cases in the test set, 52 were correctly labeled as frauds. And almost a half, 46 were mislabeled as non-frauds.
- All except 11 non-frauds were correctly labeled as non-frauds, from a total of 56,864. That's nearly perfect, but the priority should be to prevent frauds. Therefore, this is rather a secondary metric for us.

Feature Engineering

Before fixing the class imbalance, there are other things that need to be addressed:

- Classification algorithms expect to receive normalized features. There are two features in the data that aren't normalized. ('Time' and 'Amount')
- New features could be created from those unprocessed features, if they capture a pattern correlated to 'Class'.

'Features' DataFrame

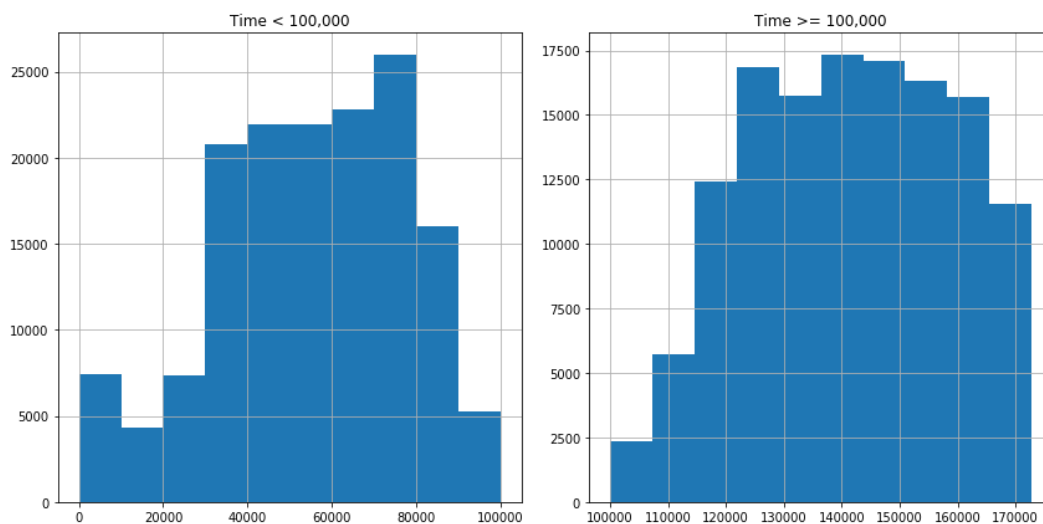
- In this dataframe I'll store the features intended for predictive modeling of frauds.
- 'data' will be left as the raw dataset.


```
1 features = pd.DataFrame()
```

Time-Based Features

- There seem to be two normal distributions in the feature Time. Let's isolate them so we can create features from them.

```
1 plt.figure(figsize=(12,6))
2
3 # Visualize where Time is less than 100,000
4 plt.subplot(1,2,1)
5 plt.title("Time < 100,000")
6 data[data['Time']<100000]['Time'].hist()
7
8 # Visualize where Time is more than 100,000
9 plt.subplot(1,2,2)
10 plt.title("Time >= 100,000")
11 data[data['Time']>=100000]['Time'].hist()
12
13 plt.tight_layout()
14 plt.show()
```




```
1 # Create a feature from normal distributions above
2 features['100k_time'] = np.where(data.Time<100000, 1,0)
```

Feature: Time_hour > 4

- Feature for non-frauds, where 'Time_hour' is above 4. This seems to have a clear differentiation.

```
1 plt.figure(figsize=(12,6))
2
3 plt.subplot(1,2,1)
4 plt.title("Non-Frauds, Hour <= 4")
5 eda.Time_hour[(eda.Class == 0) & (eda.Time_hour <= 4)].plot(kind='hist',bins=15)
6
7 plt.subplot(1,2,2)
8 plt.title("Non-Frauds, Hour > 4")
9 eda.Time_hour[(eda.Class == 0) & (eda.Time_hour > 4)].plot(kind='hist',bins=15)
10
11 plt.tight_layout()
12 plt.show()
```



```
1 # Create a feature from distributions above
2 features['4_hour'] = np.where((eda.Class == 0) & (eda.Time_hour > 4), 1,0)
```

Feature: \$0 Fraud Amounts...?

- Many transactions are zero dollars. This might be confusing for our model's predictive ability. It is arguable these don't need to be prevented.
 - One approach could be to simply discard these transactions.
 - The second approach is to ignore it and focus on predicting transactions labeled as 'frauds', regardless of them having no dollar-value.

For now, I'll use this as basis for a feature. Later I'll compare results between different approaches

```
1 # how many frauds are actually 0 dollars?
2 print("Non-Fraud Zero dollar Transactions:")
3 display(data[(data.Amount == 0) & (data.Class == 0)]['Class'].count())
4 print("Fraudulent Zero dollar Transactions:")
5 display(data[(data.Amount == 0) & (data.Class == 1)]['Class'].count())
```

```
1 Non-Fraud Zero dollar Transactions:
2 1798
3 Fraudulent Zero dollar Transactions:
4 27
```

```
1 # Capture where transactions have a $0 amount
2 features['amount0'] = np.where(data.Amount == 0,1,0)
```

Normalize Time and Amount

- Although we already captured some features from 'Time' and 'Amount', before decidedly dropping them, I'd like to normalize and test them in the model.

```
1 rob_scaler = RobustScaler()
2
3 features['scaled_amount'] = rob_scaler.fit_transform(data['Amount'].values.reshape(-1,1))
4 features['scaled_time'] = rob_scaler.fit_transform(data['Time'].values.reshape(-1,1))
```

Add the Rest: PCA and Class

```
1 # Add the PCA components to our features DataFrame.
2 features = features.join(data.iloc[:,1:-1].drop('Amount',axis=1))
3
4 # Add 'Class' to our features DataFrame.
5 features = features.join(data.Class)
6
7 # Nice! These are the final features I'll settle for.
8 features.head()
```

	100k_time	4_hour	amount0	scaled_amount	scaled_time	V1	V2	V3	V4	V5	...
0	1	0	0	1.783274	-0.994983	-1.359807	-0.072781	2.536347	1.378155	-0.338321	...
1	1	0	0	-0.269825	-0.994983	1.191857	0.266151	0.166480	0.448154	0.060018	...
2	1	0	0	4.983721	-0.994972	-1.358354	-1.340163	1.773209	0.379780	-0.503198	...
3	1	0	0	1.418291	-0.994972	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	...
4	1	0	0	0.670579	-0.994960	-1.158233	0.877737	1.548718	0.403034	-0.407193	...

5 rows × 34 columns

Classification Improvements after Feature Engineering

- We've added some features, and re-coded two existing features. Let's see how classification performs now.
- In this classification I'll define `x` and `y`, as well as `train` and `test` samples from the `features` DataFrame, which has the feature-engineered version of the data.
- Also, I'll use `recall_score` as the scoring function for cross-validation. This represents the percentage of frauds correctly identified.

```

1 # Define outcome and predictors USE FEATURE-ENGINEERED DATA
2 y = features['Class']
3 X = features.drop('Class', 1)
4
5 # Split X and y into train and test sets.
6 X_train, X_test, y_train, y_test = train_test_split(
7     X, y, test_size=0.2, random_state=42)
8
9 # Class balance in test group
10 print("TEST GROUP")
11 print('Size:', y_test.count())
12 print("Frauds percentage:",
13       y_test.value_counts()[1]/y_test.count())
14 print("Nonfrauds percentage:",
15       y_test.value_counts()[0]/y_test.count())
16
17 # Class balance in train group
18 print("\nTRAIN GROUP")
19 print('Size:', y_train.count())
20 print("Frauds percentage:",
21       y_train.value_counts()[1]/y_train.count())
22 print("Nonfrauds percentage:",
23       y_train.value_counts()[0]/y_train.count())

```

```

1 TEST GROUP
2 Size: 56962
3 Frauds percentage: 0.0017204452090867595
4 Nonfrauds percentage: 0.9982795547909132
5
6 TRAIN GROUP
7 Size: 227845
8 Frauds percentage: 0.001729245759178389
9 Nonfrauds percentage: 0.9982707542408216

```

```

1 # Invoke classifier
2 clf = LogisticRegression()
3
4 # Make a scoring callable from recall_score
5 recall = make_scorer(recall_score)
6
7 # Cross-validate on the train data
8 train_cv = cross_val_score(X=X_train, y=y_train, estimator=clf, scoring=recall, cv=3)
9 print("TRAIN GROUP")
10 print("\nCross-validation recall scores:", train_cv)
11 print("Mean recall score:", train_cv.mean())
12
13 # Now predict on the test group
14 print("\nTEST GROUP")
15 y_pred = clf.fit(X_train, y_train).predict(X_test)
16 print("\nRecall:", recall_score(y_test, y_pred))
17
18 # Classification report
19 print('\nClassification report:\n')
20 print(classification_report(y_test, y_pred))

```

```


21
22 # Confusion matrix
23 conf_matrix = confusion_matrix(y_test,y_pred)
24 sns.heatmap(conf_matrix, annot=True,fmt='d', cmap=plt.cm.copper)
25 plt.show()

```

```

1 TRAIN GROUP
2
3 Cross-validation recall scores: [0.79545455 0.81679389 0.85496183]
4 Mean recall score: 0.8224034235484617
5
6 TEST GROUP
7
8 Recall: 0.826530612244898
9
10 Classification report:
11
12           precision    recall  f1-score   support
13
14      0           1.00      1.00      1.00     56864
15      1           0.98      0.83      0.90         98
16
17   micro avg           1.00      1.00      1.00     56962
18   macro avg           0.99      0.91      0.95     56962
19   weighted avg          1.00      1.00      1.00     56962

```



Scores

- Now the cross_val scores reflect the fraud recall on three folds of the train data. These numbers are more informative for us now.
- The mean recall from train data is also very consistent with the test recall. This is evidence of the model's certainty.
- Fraud Recall went up from 53% to 83%. That's pretty good already, but it's far from perfect. We still have 17 frauds in the test set that aren't being predicted.

What's next The main obstacles for high accuracy are currently class-imbalance, outliers and noise. Fixing these involves changing the length of the data, meaning we won't have the same datapoints present afterwards. For that reason, we'll only use the features' `train` data to make these transformations, and use the features' `test` data to make predictions.

Data Processing

Data processing will include class-balancing, removing outliers, and feature-selection.

Balancing Classes

There's several methods for balancing classes: Im mostly interested in these...

- Random-Undersampling of Majority Class.

You reduce the size of majority class to match size of minority class. Disadvantage is that you may end up with very little data.

- SMOTE- Synthetic Minority Oversampling Technique.

Algorithm that creates a larger sample of minority class to match the size of majority class.

- Inverting Class Ratios. (Turning minority into majority)

If you turn the minority into the majority, you may skew results towards better recall scores(detecting frauds correctly), as opposed to better specificity scores.(detecting non-frauds correctly)

For now, I'll balance with a variant implementation of SMOTE, to see correlations.


```
1 # Balancing Classes before checking for correlation
2
3 # Join the train data
4 train = X_train.join(y_train)
5
6 print('Data shape before balancing:', train.shape)
7 print('\nCounts of frauds VS non-frauds in previous data:')
8 print(train.Class.value_counts())
9 print('-'*40)
10
11 # Oversample frauds. Imblearn's ADASYN was built for class-imbalanced datasets
12 X_bal, y_bal = ADASYN(sampling_strategy='minority', random_state=0).fit_resample(
13     X_train,
14     y_train)
15
16 # Join X and y
17 X_bal = pd.DataFrame(X_bal, columns=X_train.columns)
18 y_bal = pd.DataFrame(y_bal, columns=['Class'])
19 balanced = X_bal.join(y_bal)
20
21
22 print('-'*40)
23 print('Data shape after balancing:', balanced.shape)
24 print('\nCounts of frauds VS non-frauds in new data:')
25 print(balanced.Class.value_counts())
```

```
1 Data shape before balancing: (227845, 34)
2
3 Counts of frauds VS non-frauds in previous data:
4 0    227451
5 1     394
6 Name: Class, dtype: int64
7 -----
8 -----
9 Data shape after balancing: (454905, 34)
10
11 Counts of frauds VS non-frauds in new data:
12 1    227454
13 0    227451
14 Name: Class, dtype: int64
```

- Now we have much more data because the frauds were oversampled to match the size of non-frauds.
- Notice that ADASYN isn't perfectly matching the number of frauds to the majority class. This is good enough though.

```
1 print('Distribution of the Classes in the subsample dataset')
2 print(balanced.Class.value_counts()/len(train))
3
4 sns.countplot('Class', data=balanced)
5 plt.title('Class Distribution', fontsize=14)
6 plt.show()
7
```


```
1 Distribution of the Classes in the subsample dataset
2 1    0.998284
3 0    0.998271
4 Name: Class, dtype: float64
```



Removing High-Correlation Outliers

- This step must be taken after balancing classes. Otherwise, correlations will echo class-distributions. To illustrate, I'll include two versions of the correlation matrix.
- Based on a correlation matrix, we'll identify features with high correlations, and remove any transactions with outlying values in these.
- High correlation features have a high capacity to influence the algorithm prediction. Therefore it's important to control their anomalies.
- This approach will reduce prediction bias because our algorithm will learn from more normally-distributed features.

```
1 # Compare correlation of raw train data VS balanced train data
2
3 f, (ax1, ax2) = plt.subplots(2, 1, figsize=(24,20))
4
5 # Imbalanced DataFrame
6 corr = train.corr()
7 sns.heatmap(corr, annot_kws={'size':20}, ax=ax1)
8 ax1.set_title("Imbalanced Correlation Matrix \n (Biased)", fontsize=14)
9
10 # Balanced DataFrame
11 bal_corr = balanced.corr()
12 sns.heatmap(bal_corr, annot_kws={'size':20}, ax=ax2)
13 ax2.set_title('Balanced Correlation Matrix', fontsize=14)
14 plt.show()
```



- From the feature engineered variables, it looks like `4_hour` has a very strong (negative) correlation with 'Class'. Well, at least one was useful.
- Now let's see some actual numbers for feature correlations.

```
1 # Each feature's correlation with Class
2 bal_corr.Class
```

```
1 100k_time    0.123611
2 4_hour      -0.929016
3 amount0      0.086001
```

```

4 | scaled_amount    0.096184
5 | scaled_time      -0.121993
6 | V1               -0.231585
7 | V2               0.234517
8 | V3               -0.345542
9 | V4               0.603588
10 | V5               -0.082716
11 | V6               -0.210883
12 | V7               -0.235373
13 | V8               -0.051772
14 | V9               -0.278735
15 | V10              -0.397665
16 | V11              0.403567
17 | V12              -0.424532
18 | V13              -0.070768
19 | V14              -0.541743
20 | V15              -0.046659
21 | V16              -0.227418
22 | V17              -0.168376
23 | V18              -0.086192
24 | V19              -0.046249
25 | V20              0.019178
26 | V21              0.145141
27 | V22              -0.097711
28 | V23              -0.053527
29 | V24              -0.115729
30 | V25              0.060882
31 | V26              -0.064108
32 | V27              0.193223
33 | V28              0.127031
34 | Class            1.000000
35 | Name: Class, dtype: float64

```

- I'll make a loop that checks each feature for correlation value, and if greater than that, it'll remove outliers for that variable following a certain cutoff.

Approach to removing outliers:

For features of high positive correlation... Remove non-fraud outliers on the top range, (improve recall) and remove fraud outliers on the bottom range. (improve specificity)

For features of high negative correlation... Remove non-fraud outliers on the bottom range, (improve recall) and remove fraud outliers on the top range. (improve specificity)

```
1 | no_outliers=pd.DataFrame(balanced.copy())
```

```

1 | # Removing Outliers from high-correlation features
2 |
3 | cols = bal_corr.Class.index[:-1]
4 |
5 | # For each feature correlated with Class...
6 | for col in cols:
7 |     # If absolute correlation value is more than X percent...
8 |     correlation = bal_corr.loc['Class',col]
9 |     if np.absolute(correlation) > 0.1:
10 |
11 |         # Separate the classes of the high-correlation column
12 |         nonfrauds = no_outliers.loc[no_outliers.Class==0,col]
13 |         frauds = no_outliers.loc[no_outliers.Class==1,col]
14 |
15 |         # Identify the 25th and 75th quartiles
16 |         all_values = no_outliers.loc[:,col]
17 |         q25, q75 = np.percentile(all_values, 25), np.percentile(all_values, 75)
18 |         # Get the inter quartile range
19 |         iqr = q75 - q25
20 |         # Smaller cutoffs will remove more outliers
21 |         cutoff = iqr * 7
22 |         # Set the bounds of the desired portion to keep
23 |         lower, upper = q25 - cutoff, q75 + cutoff
24 |
25 |         # If positively correlated...
26 |         # Drop nonfrauds above upper bound, and frauds below lower bound
27 |         if correlation > 0:
28 |             no_outliers.drop(index=nonfrauds[nonfrauds>upper].index,inplace=True)
29 |             no_outliers.drop(index=frauds[frauds<lower].index,inplace=True)
30 |
31 |         # If negatively correlated...
32 |         # Drop nonfrauds below lower bound, and frauds above upper bound
33 |         elif correlation < 0:
34 |             no_outliers.drop(index=nonfrauds[nonfrauds<lower].index,inplace=True)

```

```

35     no_outliers.drop(index=frauds[frauds>upper].index,inplace=True)
36
37     print('\nData shape before removing outliers:', balanced.shape)
38     print('\nCounts of frauds VS non-frauds in previous data:')
39     print(balanced.Class.value_counts())
40     print('-'*40)
41     print('-'*40)
42     print('\nData shape after removing outliers:', no_outliers.shape)
43     print('\nCounts of frauds VS non-frauds in new data:')
44     print(no_outliers.Class.value_counts())

```

```

1  Data shape before removing outliers: (454905, 34)
2
3  Counts of frauds VS non-frauds in previous data:
4      1    227454
5      0    227451
6  Name: Class, dtype: int64
7  -----
8  -----
9
10 Data shape after removing outliers: (445647, 34)
11
12 Counts of frauds VS non-frauds in new data:
13      0    225209
14      1    220438
15  Name: Class, dtype: int64


```

- Outliers from high-correlation features are now gone. However, this created a class-imbalance again.
- I will balance the classes later when I reduce the model size. Reduction is important because classifiers may lag on high-dimensional datasets.

```

1  no_outliers.iloc[:, :-1].boxplot(rot=90,figsize=(16,4))
2  plt.title('Distributions with Less Outliers', fontsize=17)
3  plt.show()

```



Feature Selection

- I'll use the correlation matrix again, but this time I'll filter out features with low predictive power, instead of outliers.

But first, let's see what the outlier removal did to the correlations.

```

1  feat_sel = pd.DataFrame(no_outliers.copy())

2  # Make a dataframe with the class-correlations before removing outliers
3  corr_change = pd.DataFrame()
4  corr_change['correlation'] = bal_corr.Class
5  corr_change['origin'] = 'w/outliers'
6
7  # Make a dataframe with class-correlations after removing outliers
8  corr_other = pd.DataFrame()
9  corr_other['correlation'] = feat_sel.corr().Class
10 corr_other['origin'] = 'no_outliers'
11
12 # Join them
13 corr_change = corr_change.append(corr_other)
14
15 plt.figure(figsize=(14,6))
16 plt.xticks(rotation=90)
17
18 # Plot them
19 sns.set_style('darkgrid')


```



```

19 plt.title('Class Correlation per Feature. With VS W/out Outliers', fontsize=17)
20 sns.barplot(data=corr_change, x=corr_change.index, y='correlation', hue='origin')
21 plt.show()

```



- It's obvious that most features gained correlation power, regardless of direction. Positive correlations went higher up, negative correlations went lower down. Also, the highest correlations flattened out, while the smallest ones rose to relevance.
- It is clearly an indicator that the outliers were causing noise, and therefore dimming the correlation-potential of each feature.

```

1 # Feature Selection based on correlation with Class
2
3 print('\nData shape before feature selection:', feat_sel.shape)
4 print('\nCounts of frauds VS non-frauds before feature selection:')
5 print(feat_sel.Class.value_counts())
6 print('-'*40)
7
8 # Correlation matrix after removing outliers
9 new_corr = feat_sel.corr()
10
11 for col in new_corr.Class.index[:-1]:
12     # Pick desired cutoff for dropping features. In absolute-value terms.
13     if np.absolute(new_corr.loc['Class',col]) < 0.1:
14         # Drop the feature if correlation is below cutoff
15         feat_sel.drop(columns=col,inplace=True)
16
17 print('-'*40)
18 print('\nData shape after feature selection:', feat_sel.shape)
19 print('\nCounts of frauds VS non-frauds in new data:')
20 print(feat_sel.Class.value_counts())

```

```


1 Data shape before feature selection: (445647, 34)
2
3 Counts of frauds VS non-frauds before feature selection:
4 0    225209
5 1    220438
6 Name: Class, dtype: int64
7 -----
8 -----
9
10 Data shape after feature selection: (445647, 23)
11
12 Counts of frauds VS non-frauds in new data:
13 0    225209
14 1    220438
15 Name: Class, dtype: int64

```

```

1 feat_sel.iloc[:, :-1].boxplot(rot=90,figsize=(16,4))
2 plt.title('Distribution of Features Selected', fontsize=17)
3 plt.show()

```



- So this removed a few features from our 'processed' dataset. Aside from its large size, it should be ready for predictions.

Test and Compare Classifiers

Approach:

- I'll evaluate improvements based on **fraud recall**, since its crucial to prevent frauds. This might come at the expense of more false-alarms, which would decrease the overall accuracy. **The main purpose of this project will be to identify all frauds, while minimizing false-positives.**
- I'll define outcomes and predictors, reduce model size, and classify.

```
1 # Undersample model for efficiency and balance classes.
2
3 X_train = feat_sel.drop('Class',1)
4 y_train = feat_sel.Class
5
6 # After feature-selection, X_test needs to include only the same features as X_train
7 cols = X_train.columns
8 X_test = X_test[cols]
9
10 # Undersample and balance classes
11 X_train, y_train = RandomUnderSampler(sampling_strategy={1:5000,0:5000}).fit_resample(X_train,y_train)
12
13 print('\nX_train shape after reduction:', X_train.shape)
14 print('\nCounts of frauds VS non-frauds in y_train:')
15 print(np.unique(y_train, return_counts=True))
```

```
1 X_train shape after reduction: (10000, 22)
2
3 Counts of frauds VS non-frauds in y_train:
4 (array([0, 1]), array([5000, 5000]))
```

First-Run: Predictions on Default Parameters

- Here, I'll try a few simple classifiers and compare their performance.

```
1 # DataFrame to store classifier performance
2 performance = pd.DataFrame(columns=['Train_Recall','Test_Recall','Test_Specificity'])
```

```
1 # Load simple classifiers
2 classifiers = [SVC(max_iter=1000),LogisticRegression(),
3               DecisionTreeClassifier(),KNeighborsClassifier()]
4
5 # Get a classification report from each algorithm
6 for clf in classifiers:
7
8     # Heading
9     print('\n','-'*40,'\n',clf.__class__.__name__,'\n','-'*40)
10
11     # Cross-validate on the train data
12     print("TRAIN GROUP")
13     train_cv = cross_val_score(X=X_train, y=y_train,
14                               estimator=clf, scoring=recall,cv=3)
15     print("\nCross-validation recall scores:",train_cv)
16     print("Mean recall score:",train_cv.mean())
17
18     # Now predict on the test group
19     print("\nTEST GROUP")
20     y_pred = clf.fit(X_train, y_train).predict(X_test)
21     print("\nRecall:",recall_score(y_test,y_pred))
22
```

```


23 # Print confusion matrix
24 conf_matrix = confusion_matrix(y_test,y_pred)
25 sns.heatmap(conf_matrix, annot=True,fmt='d', cmap=plt.cm.copper)
26 plt.show()
27
28 # Store results
29 performance.loc[clf.__class__.__name__+'_default',
30                ['Train_Recall','Test_Recall','Test_Specificity']] = [
31    train_cv.mean(),
32    recall_score(y_test,y_pred),
33    conf_matrix[0,0]/conf_matrix[0,:].sum()
34 ]

```

```

1  -----
2  SVC
3  -----
4  TRAIN GROUP
5
6  Cross-validation recall scores: [0.99940012 1.          1.          ]
7  Mean recall score: 0.9998000399920016
8
9  TEST GROUP
10
11 Recall: 0.7653061224489796

```



```


----- LogisticRegression ----- TRAIN GROUP
Cross-validation recall scores: [0.99880024 0.9970006 0.99759904] Mean recall score: 0.99779995981596

```

```

1  TEST GROUP
2
3  Recall: 0.9897959183673469

```



```


----- DecisionTreeClassifier ----- TRAIN GROUP
Cross-validation recall scores: [1. 0.99760048 0.99639856] Mean recall score: 0.9979996797759295

```

```


1  TEST GROUP
2
3  Recall: 0.9897959183673469

```



----- KNeighborsClassifier ----- TRAIN GROUP
 Cross-validation recall scores: [1. 0.99940012 1.] Mean recall score: 0.9998000399920016

```
1 TEST GROUP
2
3 Recall: 0.9081632653061225
```



```
1 # Scores obtained
2 performance
```

	Train_Recall	Test_Recall	Test_Specificity
SVC_default	0.9998	0.765306	0.991119
LogisticRegression_default	0.9978	0.989796	0.976611
DecisionTreeClassifier_default	0.998	0.989796	0.993739
KNeighborsClassifier_default	0.9998	0.908163	0.96722

- These results are very promising for a first run, considering I didn't tweak any of the parameters.
- Now let's do a GridSearchCV to find the best parameters for these classifiers.

Logistic Regression- GridSearch & Recall Score.

- `GridSearchCV` compares parameter combinations to find the highest score, determined by the user. I'll set `recall_score` to be the determinant factor for the best parameter combination.
- The `class_weight` parameter greatly skews the classification emphasis from focusing on frauds at the expense of more non-fraud errors. For now, I'll prioritize fraud prevention. Later, I'll improve on specificity.

About the parameters to optimize

- Solvers `'newton-cg'`, `'lbfgs'`, and `'sag'` handle only `L2`-penalty. So we'll have to do this using two parameter grids: First for `L2`-only solvers, and then for `L1` and `L2`-solvers.

```
1 # Parameters to optimize
2 params = [{
3     'solver': ['newton-cg', 'lbfgs', 'sag'],
4     'C': [0.3, 0.5, 0.7, 1],
5     'penalty': ['l2']
6 }, {
7     'solver': ['liblinear', 'saga'],
8     'C': [0.3, 0.5, 0.7, 1],
9     'penalty': ['l1', 'l2']
10 }]
```

```


10  }}
11
12  clf = LogisticRegression(
13      n_jobs=-1, # Use all CPU
14      class_weight={0:0.1,1:1} # Prioritize frauds
15  )
16
17  # Load GridSearchCV
18  search = GridSearchCV(
19      estimator=clf,
20      param_grid=params,
21      n_jobs=-1,
22      scoring=recall
23  )
24
25  # Train search object
26  search.fit(X_train, y_train)
27
28  # Heading
29  print('\n','-'*40,'\n',clf.__class__.__name__,'\n','-'*40)
30
31  # Extract best estimator
32  best = search.best_estimator_
33  print('Best parameters: \n\n',search.best_params_,'\n')
34
35  # Cross-validate on the train data
36  print("TRAIN GROUP")
37  train_cv = cross_val_score(X=X_train, y=y_train,
38                             estimator=best, scoring=recall,cv=3)
39  print("\nCross-validation recall scores:",train_cv)
40  print("Mean recall score:",train_cv.mean())
41
42  # Now predict on the test group
43  print("\nTEST GROUP")
44  y_pred = best.fit(X_train, y_train).predict(X_test)
45  print("\nRecall:",recall_score(y_test,y_pred))
46
47  # Get classification report
48  print(classification_report(y_test, y_pred))
49
50  # Print confusion matrix
51  conf_matrix = confusion_matrix(y_test,y_pred)
52  sns.heatmap(conf_matrix, annot=True, fmt='d', cmap=plt.cm.copper)
53  plt.show()
54
55  # Store results
56  performance.loc[clf.__class__.__name__+'_search',
57                 ['Train_Recall','Test_Recall','Test_Specificity']] = [
58      train_cv.mean(),
59      recall_score(y_test,y_pred),
60      conf_matrix[0,0]/conf_matrix[0,:].sum()
61  ]

```

```

1  -----
2  LogisticRegression
3  -----
4  Best parameters:
5
6  {'C': 0.3, 'penalty': 'l2', 'solver': 'newton-cg'}
7
8  TRAIN GROUP
9
10 Cross-validation recall scores: [1. 1. 1.]
11 Mean recall score: 1.0
12
13 TEST GROUP
14
15 Recall: 1.0
16
17      precision    recall  f1-score   support
18
19  0           1.00      0.93      0.97       56864
20  1           0.03      1.00      0.05          98
21
22  micro avg       0.93      0.93      0.93       56962
23  macro avg       0.51      0.97      0.51       56962
24  weighted avg     1.00      0.93      0.96       56962

```



```
1 | performance
```

	Train_Recall	Test_Recall	Test_Specificity
SVC_default	0.9998	0.765306	0.991119
LogisticRegression_default	0.9978	0.989796	0.976611
DecisionTreeClassifier_default	0.998	0.989796	0.993739
KNeighborsClassifier_default	0.9998	0.908163	0.96722
LogisticRegression_search	1	1	0.934018

Pyrrhic Victory-

A victory that inflicts such a devastating toll on the victor that it is tantamount to defeat. Someone who wins a Pyrrhic victory has also taken a heavy toll that negates any true sense of achievement.

- Well, fraud recall improved on Logistic Regression.
- However, this has come at the cost of horribly low specificity.
- `GridSearch` allows us to see the results that informed the choice of best parameters, based on our scoring function. In this case, `recall_score`. Let's see how they compare.

```
1 | pd.DataFrame(search.cv_results_).iloc[:,4:].sort_values(by='rank_test_score').head()
```

	param_C	param_penalty	param_solver	params	split0_test_score	split1_test_score	split2_test_score	mean_test_score
0	0.3	l2	newton-cg	{'C': 0.3, 'penalty': 'l2', 'solver': 'newton-...	1.0	1.0	1.0	1.0
25	1	l1	saga	{'C': 1, 'penalty': 'l1', 'solver': 'saga'}	1.0	1.0	1.0	1.0
24	1	l1	liblinear	{'C': 1, 'penalty': 'l1', 'solver': 'liblinear'}	1.0	1.0	1.0	1.0
23	0.7	l2	saga	{'C': 0.7, 'penalty': 'l2', 'solver': 'saga'}	1.0	1.0	1.0	1.0
22	0.7	l2	liblinear	{'C': 0.7, 'penalty': 'l2', 'solver': 'libline...	1.0	1.0	1.0	1.0

- It seems like the top 5 combinations had a perfect `recall_score`, which explain why they all have a rank of `1`. This means there was no need for a real comparison for the 'best' parameters, because they all were perfect. We simply got the

parameters that were first on the list of **perfect** combinations.

- Since we wanted to prioritize fraud recall, we set a very skewed `class_weight` parameter. This is why the results produced such perfect recall scores, at the expense of specificity.
- Let's find the right balance between perfect recall and higher specificity.

Optimize Specificity, while Maintaining 100% Recall

- In this section I'll implement a few ideas to minimize false-positives (non-frauds identified as frauds), while still predicting all frauds correctly.

Custom Scoring Function

- Parameter search functions use a scoring parameter to determine the best parameter combination. In the previous experiments we've used recall score as the basis. Now we want to pick a parameter combination that also takes specificity into account, while ensuring perfect recall.

```
1 # Make a scoring function that improves specificity while identifying all frauds
2 def recall_optim(y_true, y_pred):
3
4     conf_matrix = confusion_matrix(y_true, y_pred)
5
6     # Recall will be worth a greater value than specificity
7     rec = recall_score(y_true, y_pred) * 0.8
8     spe = conf_matrix[0,0]/conf_matrix[0,:].sum() * 0.2
9
10    # Imperfect recalls will lose a penalty
11    # This means the best results will have perfect recalls and compete for specificity
12    if rec < 0.8:
13        rec -= 0.2
14    return rec + spe
15
16 # Create a scoring callable based on the scoring function
17 optimize = make_scorer(recall_optim)
```

- Now add the optimized scores to the existing performance DataFrame

```
1 scores = []
2 for rec, spe in performance[['Test_Recall', 'Test_Specificity']].values:
3     rec = rec * 0.8
4     spe = spe * 0.2
5     if rec < 0.8:
6         rec -= 0.20
7     scores.append(rec + spe)
8 performance['Optimize'] = scores
9 display(performance)
```

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804

Iteration Function

Since I'll apply the new settings to several classifiers, I'll define a function to reuse several times.

- It'll take the parameters you want to compare, and the classifier you want to try.
- It'll determine best parameters based on custom scoring, do cross-validation for recall on train data, then train and predict the test set.
- It'll show us the recall scores for train and test, a confusion matrix, a classification report, the GridSearch' top combinations, and a view of the performance DataFrame.

```
1 def score_optimization(params,clf):
2     # Load GridSearchCV
3     search = GridSearchCV(
4         estimator=clf,
5         param_grid=params,
6         n_jobs=-1,
7         scoring=optimize
8     )
9
```

```

10 # Train search object
11 search.fit(X_train, y_train)
12
13 # Heading
14 print('\n', '-'*40, '\n', clf.__class__.__name__, '\n', '-'*40)
15
16 # Extract best estimator
17 best = search.best_estimator_
18 print('Best parameters: \n\n', search.best_params_, '\n')
19
20 # Cross-validate on the train data
21 print("TRAIN GROUP")
22 train_cv = cross_val_score(X=X_train, y=y_train,
23                             estimator=best, scoring=recall, cv=3)
24 print("\nCross-validation recall scores:", train_cv)
25 print("Mean recall score:", train_cv.mean())
26
27 # Now predict on the test group
28 print("\nTEST GROUP")
29 y_pred = best.fit(X_train, y_train).predict(X_test)
30 print("\nRecall:", recall_score(y_test, y_pred))
31
32 # Get classification report
33 print(classification_report(y_test, y_pred))
34
35 # Print confusion matrix
36 conf_matrix = confusion_matrix(y_test, y_pred)
37 sns.heatmap(conf_matrix, annot=True, fmt='d', cmap=plt.cm.copper)
38 plt.show()
39
40 # Store results
41 performance.loc[clf.__class__.__name__+'_optimize',:] = [
42     train_cv.mean(),
43     recall_score(y_test, y_pred),
44     conf_matrix[0,0]/conf_matrix[0,:].sum(),
45     recall_optim(y_test, y_pred)
46 ]
47 # Look at the parameters for the top best scores
48 display(pd.DataFrame(search.cv_results_).iloc[:,4:].sort_values(by='rank_test_score').head())
49 display(performance)

```

LogisticRegression- Optimized.

```

1 # Parameters to optimize
2 params = [{
3     'solver': ['newton-cg', 'lbfgs', 'sag'],
4     'C': [0.3, 0.5, 0.7, 1],
5     'penalty': ['l2'],
6     'class_weight': [{1:1,0:0.3},{1:1,0:0.5},{1:1,0:0.7}]
7 }, {
8     'solver': ['liblinear', 'saga'],
9     'C': [0.3, 0.5, 0.7, 1],
10    'penalty': ['l1', 'l2'],
11    'class_weight': [{1:1,0:0.3},{1:1,0:0.5},{1:1,0:0.7}]
12 }]
13
14 clf = LogisticRegression(
15     n_jobs=-1 # Use all CPU
16 )
17
18 score_optimization(clf=clf, params=params)


```

```

1 -----
2 LogisticRegression
3 -----
4 Best parameters:
5
6 {'C': 1, 'class_weight': {1: 1, 0: 0.5}, 'penalty': 'l1', 'solver': 'liblinear'}
7
8 TRAIN GROUP
9
10 Cross-validation recall scores: [1. 1. 1.]
11 Mean recall score: 1.0
12
13 TEST GROUP
14
15 Recall: 1.0
16
17 precision    recall  f1-score   support

```


18	0	1.00	0.98	0.99	56864
19	1	0.07	1.00	0.13	98
20					
21	micro avg	0.98	0.98	0.98	56962
22	macro avg	0.53	0.99	0.56	56962
23	weighted avg	1.00	0.98	0.99	56962



	param_C	param_class_weight	param_penalty	param_solver	params	split0_test_score	split1_test_score	split2_test_s
76	1	{1: 1, 0: 0.5}	l1	liblinear	{'C': 1, 'class_weight': {1: 1, 0: 0.5}, 'pena...	0.994481	0.995561	0.995558
64	0.7	{1: 1, 0: 0.5}	l1	liblinear	{'C': 0.7, 'class_weight': {1: 1, 0: 0.5}, 'pe...	0.994601	0.995201	0.995438
82	1	{1: 1, 0: 0.7}	l2	liblinear	{'C': 1, 'class_weight': {1: 1, 0: 0.7}, 'pena...	0.994481	0.994241	0.995438
52	0.5	{1: 1, 0: 0.5}	l1	liblinear	{'C': 0.5, 'class_weight': {1: 1, 0: 0.5}, 'pe...	0.994241	0.994841	0.995078
70	0.7	{1: 1, 0: 0.7}	l2	liblinear	{'C': 0.7, 'class_weight': {1: 1, 0: 0.7}, 'pe...	0.994481	0.994001	0.995198


	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357

- Yes!! With our optimize function, specificity in LogisticRegression improved from 93% to 97%, while still having perfect recall.
- By looking at these results, there's no doubt that `liblinear`, `l1` is the best combination, regardless of `C_param`.
- Also, `class_weight` for non-frauds set to 0.5 (`1:1,0:5`) seem to rank better. This is likely the result of the custom scoring which now rewards higher precisions.

DecisionTreeClassifier- Optimized

```
1 # Parameters to optimize
2 params = {
3     'criterion':['gini','entropy'],
4     'max_features':[None,'sqrt'],
5     'class_weight':[{1:1,0:0.3},{1:1,0:0.5},{1:1,0:0.7}]
6 }
7
8 clf = DecisionTreeClassifier(
9 )
10
11 score_optimization(clf=clf,params=params)
```

```
1 -----
2 DecisionTreeClassifier
3 -----
4 Best parameters:
5
6 {'class_weight': {1: 1, 0: 0.5}, 'criterion': 'gini', 'max_features': None}
7
8 TRAIN GROUP
9
10 Cross-validation recall scores: [0.99640072 0.99640072 0.99759904]
11 Mean recall score: 0.9968001597759679
12
13 TEST GROUP
14
15 Recall: 0.9693877551020408
16 precision    recall  f1-score   support
17
18      0       1.00      0.99      1.00     56864
19      1       0.18      0.97      0.30        98
20
21   micro avg       0.99       0.99       0.99     56962
22   macro avg       0.59       0.98       0.65     56962
23  weighted avg       1.00       0.99       0.99     56962
```



	param_class_weight	param_criterion	param_max_features	params	split0_test_score	split1_test_score	split2_test_score
4	{1: 1, 0: 0.5}	gini	None	{'class_weight': {1: 1, 0: 0.5}, 'criterion': ...}	0.797001	0.796161	0.796519
8	{1: 1, 0: 0.7}	gini	None	{'class_weight': {1: 1, 0: 0.7}, 'criterion': ...}	0.797121	0.795321	0.796279
6	{1: 1, 0: 0.5}	entropy	None	{'class_weight': {1: 1, 0: 0.5}, 'criterion': ...}	0.795921	0.796281	0.796158
2	{1: 1, 0: 0.3}	entropy	None	{'class_weight': {1: 1, 0: 0.3}, 'criterion': ...}	0.796641	0.796761	0.794718
10	{1: 1, 0: 0.7}	entropy	None	{'class_weight': {1: 1, 0: 0.7}, 'criterion': ...}	0.797121	0.794241	0.795798

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938

- All the `None` parameters performed better.
- By looking at the top split_scores, several are less than 0.8, which means not-perfect recalls. No wonder it didn't nail all the frauds.
- DecisionTreeClassifier seems to be better at predicting non-frauds than others, but consistently misses a few frauds.
- Between default and optimize scores, DecisionTree lost accuracy. Well, some algorithms have their limitations.

Support Vector Classifier- Optimized

- The fit time complexity is more than quadratic with the number of samples which makes it hard to scale to dataset with more than a couple of 10000 samples. <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>
- C is 1 by default and it's a reasonable default choice. If you have a lot of noisy observations you should decrease it. It corresponds to regularize more the estimation. <https://scikit-learn.org/stable/modules/svm.html#tips-on-practical-use>

```


1  # Parameters to optimize
2  params = {
3      'kernel': ['rbf', 'linear'],
4      'C': [0.3, 0.5, 0.7, 1],
5      'gamma': ['auto', 'scale'],
6      'class_weight': [{1: 1, 0: 0.3}, {1: 1, 0: 0.5}, {1: 1, 0: 0.7}]
7  }
8
9  # Load classifier
10 clf = SVC(
11     cache_size=3000,
12     max_iter=1000, # Limit processing time
13 )
14 score_optimization(clf=clf, params=params)

```

```

1  -----
2  SVC
3  -----
4  Best parameters:
5
6  {'C': 0.7, 'class_weight': {1: 1, 0: 0.7}, 'gamma': 'auto', 'kernel': 'rbf'}
7
8  TRAIN GROUP
9
10 Cross-validation recall scores: [1. 1. 1.]
11 Mean recall score: 1.0
12
13 TEST GROUP
14
15 Recall: 0.7653061224489796
16
17      precision    recall  f1-score   support
18
19     0           1.00      0.99      0.99     56864
20     1           0.09      0.77      0.16         98
21
22    micro avg       0.99      0.99      0.99     56962
23    macro avg       0.55      0.88      0.58     56962
24   weighted avg       1.00      0.99      0.99     56962

```



	param_C	param_class_weight	param_gamma	param_kernel	params	split0_test_score	split1_test_score	split2_test_s
32	0.7	{1: 1, 0: 0.7}	auto	rbf	{'C': 0.7, 'class_weight': {1: 1, 0: 0.7}, 'ga...	0.996161	0.996041	0.997239
40	1	{1: 1, 0: 0.5}	auto	rbf	{'C': 1, 'class_weight': {1: 1, 0: 0.5}, 'gamm...	0.996041	0.996041	0.997239
20	0.5	{1: 1, 0: 0.7}	auto	rbf	{'C': 0.5, 'class_weight': {1: 1, 0: 0.7}, 'ga...	0.994961	0.995201	0.996639
46	1	{1: 1, 0: 0.7}	scale	rbf	{'C': 1, 'class_weight': {1: 1, 0: 0.7}, 'gamm...	0.994961	0.995441	0.996158
28	0.7	{1: 1, 0: 0.5}	auto	rbf	{'C': 0.7, 'class_weight': {1: 1, 0: 0.5}, 'ga...	0.994961	0.994841	0.996519

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938
SVC_optimize	1	0.765306	0.986811	0.609607

- SVC's scores have the most disparity between train and test sets. Train splits had perfect recall, but test set was very poor.
- First three have `param_C` to `1`, followed by `0.7`. That's very conclusive I'd say.
- Compared with its default settings, its score also decreased. SVC can be very good at learning from train data, but it's very sensitive when tested in different data.

KNeighborsClassifier- Optimized

```


1 # Parameters to compare
2 params = {
3     "n_neighbors": list(range(2,6,1)),
4     'leaf_size': list(range(20,41,10)),
5     'algorithm': ['ball_tree','auto'],
6     'p': [1,2] # Regularization parameter. Equivalent to 'l1' or 'l2'
7 }
8
9 # Load classifier
10 clf = KNeighborsClassifier(
11     n_jobs=-1
12 )
13 score_optimization(clf=clf,params=params)

```

```

1 -----
2 KNeighborsClassifier
3 -----
4 Best parameters:
5
6 {'algorithm': 'ball_tree', 'leaf_size': 20, 'n_neighbors': 2, 'p': 1}
7
8 TRAIN GROUP
9
10 Cross-validation recall scores: [1. 1. 1.]
11 Mean recall score: 1.0
12
13 TEST GROUP
14
15 Recall: 0.8979591836734694
16
17      precision    recall  f1-score   support
18
19     0       1.00      0.99      1.00     56864
20     1       0.14      0.90      0.24         98
21
22    micro avg       0.99      0.99      0.99     56962
23    macro avg       0.57      0.94      0.62     56962
24    weighted avg      1.00      0.99      0.99     56962

```



	param_algorithm	param_leaf_size	param_n_neighbors	param_p	params	split0_test_score	split1_test_score	split2_test_score
0	ball_tree	20	2	1	{'algorithm': 'ball_tree', 'leaf_size': 20, 'n_neighbors': 2, 'p': 1}	0.997361	0.99784	0.997719
40	auto	40	2	1	{'algorithm': 'auto', 'leaf_size': 40, 'n_neighbors': 2, 'p': 1}	0.997361	0.99784	0.997719
16	ball_tree	40	2	1	{'algorithm': 'ball_tree', 'leaf_size': 40, 'n_neighbors': 2, 'p': 1}	0.997361	0.99784	0.997719
24	auto	20	2	1	{'algorithm': 'auto', 'leaf_size': 20, 'n_neighbors': 2, 'p': 1}	0.997361	0.99784	0.997719
32	auto	30	2	1	{'algorithm': 'auto', 'leaf_size': 30, 'n_neighbors': 2, 'p': 1}	0.997361	0.99784	0.997719

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938
SVC_optimize	1	0.765306	0.986811	0.609607
KNeighborsClassifier_optimize	1	0.897959	0.990328	0.716433


Imblearn' BalancedRandomForest- Optimized

- This algorithm incorporates a RandomForestClassifier with a RandomUndersampling algorithm to balance classes according to the `sampling_strategy` parameter.

```
1 # Parameters to compare
2 params = {
3     'class_weight': [{1:1,0:0.3},{1:1,0:0.4},{1:1,0:0.5},{1:1,0:0.6},{1:1,0:0.7}],
4     'sampling_strategy':['all', 'not majority', 'not minority']
5 }
6
7 # Implement the classifier
8 clf = BalancedRandomForestClassifier(
9     criterion='entropy',
10    max_features=None,
11    n_jobs=-1
12 )
13 score_optimization(clf=clf,params=params)
```

```
1 -----
2 BalancedRandomForestClassifier
3 -----
4 Best parameters:
5
6 {'class_weight': {1: 1, 0: 0.6}, 'sampling_strategy': 'all'}
```

8	TRAIN GROUP				
9					
10	Cross-validation recall scores: [1.00000000 0.99940012 0.99819928]				
11	Mean recall score: 0.9991997998959632				
12					
13	TEST GROUP				
14					
15	Recall: 0.9897959183673469				
16		precision	recall	f1-score	support
17					
18	0	1.00	0.99	1.00	56864
19	1	0.22	0.99	0.36	98
20					
21	micro avg	0.99	0.99	0.99	56962
22	macro avg	0.61	0.99	0.68	56962
23	weighted avg	1.00	0.99	1.00	56962



	param_class_weight	param_sampling_strategy	params	split0_test_score	split1_test_score	split2_test_score	mean_test_score
9	{1: 1, 0: 0.6}	all	{'class_weight': {1: 1, 0: 0.6}, 'sampling_strat...	0.99832	0.99856	0.798439	0.93180
8	{1: 1, 0: 0.5}	not minority	{'class_weight': {1: 1, 0: 0.5}, 'sampling_strat...	0.99868	0.79784	0.998800	0.93176
5	{1: 1, 0: 0.4}	not minority	{'class_weight': {1: 1, 0: 0.4}, 'sampling_strat...	0.99868	0.79772	0.998800	0.93172
0	{1: 1, 0: 0.3}	all	{'class_weight': {1: 1, 0: 0.3}, 'sampling_strat...	0.99844	0.79784	0.998920	0.93172
12	{1: 1, 0: 7}	all	{'class_weight': {1: 1, 0: 7}, 'sampling_strat...	0.99952	0.99832	0.797119	0.93168

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938
SVC_optimize	1	0.765306	0.986811	0.609607
KNeighborsClassifier_optimize	1	0.897959	0.990328	0.716433
BalancedRandomForestClassifier_optimize	0.9992	0.989796	0.99395	0.790627


- Our best overall scores on test group. Recal wasn't perfect, but it has the highest combination of scores.

SKlearn' RandomForestClassifier- Optimized

- This is the good ol' RandomForestClassifier from Sklearn. It's a less specialized implementation. We'll see how it stacks against Imblearn's implementation.

```
1 # Parameters to compare
2 params = {
3     'criterion':['entropy','gini'],
4     'class_weight':[{1:1,0:0.3},{1:1,0:0.4},{1:1,0:0.5},{1:1,0:0.6},{1:1,0:0.7}]
5 }
6
7 # Implement the classifier
8 clf = RandomForestClassifier(
9     n_estimators=100,
10    max_features=None,
11    n_jobs=-1,
12 )
13
14 score_optimization(clf=clf,params=params)
```

```
1 -----
2 RandomForestClassifier
3 -----
4 Best parameters:
5
6 {'class_weight': {1: 1, 0: 0.7}, 'criterion': 'entropy'}
7
8 TRAIN GROUP
9
10 Cross-validation recall scores: [1.          1.          0.99819928]
11 Mean recall score: 0.9993997599039616
12
13 TEST GROUP
14
15 Recall: 0.9897959183673469
16
17      precision    recall  f1-score   support
18
19     0       1.00      1.00      1.00     56864
20     1       0.27      0.99      0.42        98
21
22    micro avg       1.00      1.00      1.00     56962
23    macro avg       0.63      0.99      0.71     56962
24    weighted avg       1.00      1.00      1.00     56962
```



	param_class_weight	param_criterion	params	split0_test_score	split1_test_score	split2_test_score	mean_test_score	s
8	{1: 1, 0: 7}	entropy	{'class_weight': {1: 1, 0: 7}, 'criterion': 'e...	0.99904	0.99844	0.797719	0.93176	0
9	{1: 1, 0: 7}	gini	{'class_weight': {1: 1, 0: 7}, 'criterion': 'g...	0.99868	0.99808	0.797599	0.93148	0
7	{1: 1, 0: 0.6}	gini	{'class_weight': {1: 1, 0: 0.6}, 'criterion': ...	0.99844	0.99784	0.797479	0.93128	0
3	{1: 1, 0: 0.4}	gini	{'class_weight': {1: 1, 0: 0.4}, 'criterion': ...	0.99808	0.99808	0.797359	0.93120	0
0	{1: 1, 0: 0.3}	entropy	{'class_weight': {1: 1, 0: 0.3}, 'criterion': ...	0.99844	0.79808	0.798319	0.86496	0

	Train_Recall	Test_Recall	Test_Specificity	Optimize
SVC_default	0.9998	0.765306	0.991119	0.610469
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975
LogisticRegression_search	1	1	0.934018	0.986804
LogisticRegression_optimize	1	1	0.976787	0.995357
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938
SVC_optimize	1	0.765306	0.986811	0.609607
KNeighborsClassifier_optimize	1	0.897959	0.990328	0.716433
BalancedRandomForestClassifier_optimize	0.9992	0.989796	0.99395	0.790627
RandomForestClassifier_optimize	0.9994	0.989796	0.99534	0.790905

```

1 # Let's get the mean between test recall and test specificity
2 performance['Mean_RecSpe'] = (performance.Test_Recall+performance.Test_Specificity)/2
3 performance

```

	Train_Recall	Test_Recall	Test_Specificity	Optimize	Mean_RecSpe
SVC_default	0.9998	0.765306	0.991119	0.610469	0.878213
LogisticRegression_default	0.9978	0.989796	0.976611	0.787159	0.983203
DecisionTreeClassifier_default	0.998	0.989796	0.993739	0.790585	0.991768
KNeighborsClassifier_default	0.9998	0.908163	0.96722	0.719975	0.937692
LogisticRegression_search	1	1	0.934018	0.986804	0.967009
LogisticRegression_optimize	1	1	0.976787	0.995357	0.988393
DecisionTreeClassifier_optimize	0.9968	0.969388	0.992139	0.773938	0.980763
SVC_optimize	1	0.765306	0.986811	0.609607	0.876058
KNeighborsClassifier_optimize	1	0.897959	0.990328	0.716433	0.944143
BalancedRandomForestClassifier_optimize	0.9992	0.989796	0.99395	0.790627	0.991873
RandomForestClassifier_optimize	0.9994	0.989796	0.99534	0.790905	0.992568

4. Research Question

What is the best way to predict frauds? (Pick an approach...)

- Focus on reducing false negatives.

VS

- Focus on reducing false positives.

VS

- Focus on a custom balance?

5. Choosing Model

Perfect Recall

- Judged by perfect recall and high specificity, LogisticRegression had the highest optimized score with 97% specificity and 100% recall.

Best Overall

- For a more flexible approach, RandomForestClassifier had the highest combined recall and specificity with only one missed fraud and 99% specificity.

6. Practical Use for Audiences of Interest

- **Bank's fraud-prevention mechanisms.** (Annoying: Transactions canceled when traveling)
- **Data Science students.** Addition to the pool of Kaggle's forks on this Dataset.

7. Weak Points & Shortcomings

- **Model Processing-** Involves many steps. Steps depend immensely on the data. This doesn't lend itself to quick iterations.

Could've used a processing pipeline function, but that's a more advanced method I haven't experimented with.

- **Need for Data Reduction-** 270,000 non-frauds were undersampled to 5,000... Definitely affected accuracy. A supercomputer might handle complete set without the need for reduction. SVM and KNeighbors took the longest, even after undersampling the train data.